

Freddy Shaikh

AI & AGENTIC SYSTEMS ENGINEER · SENIOR FULL STACK SOFTWARE ENGINEER

Toronto, Canada

+1(437) 267-3291 | shaikhfh1@gmail.com | freddyshaikh.com | Freddy-S3 | in freddy-shaikh

Summary

AI and agentic systems engineer with 7+ years of full-stack, cloud-native delivery on AWS underneath it. Designated technical lead for AI-driven engineering acceleration at Morningstar: architected the team’s agentic engineering harness, led MCP server architecture as technical SME, and shipped an LLM-driven audio product end-to-end on Amazon Bedrock and Polly.

Skills

AI & Agentic	LLM Orchestration, Amazon Bedrock, MCP (Model Context Protocol) Servers, Agentic Coding Harnesses, Autonomous Agent Workflows, Prompt & Context Engineering, Token Budgeting, Agent Guardrails
Languages	Python, Go, Java, C#, TypeScript, JavaScript, SQL, PowerShell, Bash, Gremlin (Graph Query)
Cloud & Platform	AWS (Lambda, ECS, S3, CloudFront, Bedrock, Polly, SNS, IAM), GCP, Docker, Kubernetes, Terraform, IaC, CI/CD
Application	Microservices, Distributed Systems, REST, GraphQL, Knowledge Graphs, .NET, Spring Boot, Vue 3, Astro, Playwright

AI & Agentic Engineering

Production AI Delivery - Morningstar

- Led MCP (Model Context Protocol) server architecture as technical SME, from proposal and technical design review through production release.
- Owned end-to-end delivery of an AI audio feature: .NET microservices on AWS Bedrock and Amazon Polly for LLM-driven SSML speech synthesis, plus the Vue 3 front end shipping transcript and smart-summary formats.
- Rolled the harness out as team practice - reusable skill files, scoped context and token budgets, operating-mode guardrails - mentoring senior and principal engineers onto it as their default workflow.

Agentic Engineering Harness - open source (GitHub)

- Provider-neutral harness for autonomous coding agents across GitHub Copilot, Claude Code, and Codex: one skill format projected into each host by an idempotent installer, operating modes separating supervised from unattended posture, and a usage-limit-aware queue that resumes after rate-limit resets.

The Compounding Engineer - writing (GitHub)

- SEO- and ad-monetized technical blog on Astro covering agentic harness engineering and self-directed learning.

Platform Experience

Morningstar

Toronto, Canada

FULL STACK SOFTWARE ENGINEER

January 2025 - PRESENT

- Built a GraphQL API over a Gremlin knowledge graph to model and query concept relationships, improving content discoverability across the platform.
- Designed Harness CI/CD pipelines with Terraform Infrastructure as Code, cutting AWS Lambda and ECS deployment time.

Santoku Corporation / Rimm.ai

Toronto, Canada

SOFTWARE ENGINEER

August 2019 - January 2025

- Led the rearchitecture of a monolithic on-premises application into AWS microservices on Docker and Kubernetes, and owned custom haptic VR training systems for enterprise clients end-to-end as primary technical decision-maker.

Japan Exchange Teaching (JET) Program

Tokyo, Japan

ENGLISH LANGUAGE TEACHER

August 2018 - August 2019

- Taught over 1,000 elementary and 500 middle school students, adapting curriculum using self-taught Japanese (JLPT N2).

Certifications & Education

2026	AWS Certified Machine Learning Engineer - Associate, AWS Training and Certification	Toronto, Canada
2025	AWS Certified AI Practitioner, AWS Training and Certification	Toronto, Canada
2023	AWS Certified Solutions Architect - Professional, AWS Training and Certification	Toronto, Canada
2018	Bachelor of Honors Science - Kinesiology, McMaster University	Hamilton, Canada