

Freddy Shaikh

AI & AGENTIC SYSTEMS ENGINEER · SENIOR FULL STACK SOFTWARE ENGINEER

Toronto, Canada

📞 +1(437) 267-3291 | ✉️ shaikhfh1@gmail.com | 🏠 freddyshaikh.com | 🎧 Freddy-S3 | in freddy-shaikh

Summary

AI and agentic systems engineer with 7+ years of full-stack, cloud-native delivery on AWS underneath it. Designated technical lead for AI-driven engineering acceleration at Morningstar: architected the team's agentic engineering harness, led MCP server architecture as technical SME, and shipped an LLM-driven audio product end-to-end on Amazon Bedrock and Polly.

Skills

AI & Agentic	LLM Orchestration, Amazon Bedrock, MCP (Model Context Protocol) Servers, Agentic Coding Harnesses, Autonomous Agent Workflows, Prompt & Context Engineering, Token Budgeting, Agent Guardrails
Languages	Python, Go, Java, C#, TypeScript, JavaScript, SQL, PowerShell, Bash, Gremlin (Graph Query)
Cloud & Platform	AWS (Lambda, ECS, S3, CloudFront, Bedrock, Polly, SNS, IAM), GCP, Docker, Kubernetes, Terraform, IaC, CI/CD
Application	Microservices, Distributed Systems, REST, GraphQL, .NET, Spring Boot, Vue 3, Playwright, Tableau

AI & Agentic Engineering

Production AI Delivery - Morningstar

- Led MCP (Model Context Protocol) server architecture as technical SME, from proposal and technical design review through production release.
- Owned end-to-end delivery of an AI audio feature: .NET microservices on AWS Bedrock and Amazon Polly for LLM-driven SSML speech synthesis, plus the Vue 3 front end shipping transcript and smart-summary formats.
- Rolled the harness out as team practice - reusable skill files, scoped context and token budgets, operating-mode guardrails - mentoring senior and principal engineers onto it as their default workflow.

Agentic Engineering Harness - open source (GitHub)

- Provider-neutral harness for autonomous coding agents across GitHub Copilot, Claude Code, and Codex: one skill format projected into each host by an idempotent installer, operating modes separating supervised from unattended posture, and a usage-limit-aware queue that resumes after rate-limit resets.
- Concurrency and recovery model for parallel agents: exclusive working-tree claims with heartbeats and stale-claim takeover, git-worktree isolation, and compare-and-swap writes to shared state - enforced by a pre-commit hook that rejects an unclaimed agent write, with task state journaled before each risky step so an interrupted run resumes deterministically.

Platform Experience

Morningstar

Toronto, Canada

FULL STACK SOFTWARE ENGINEER

January 2025 - PRESENT

- Built a GraphQL API over a Gremlin knowledge graph to model and query concept relationships, improving content discoverability across the platform.
- Designed Harness CI/CD pipelines in YAML for AWS Lambda and ECS deploys, cutting deployment time.

Santoku Corporation / Rimm.ai

Toronto, Canada

SOFTWARE ENGINEER

August 2019 - January 2025

- Led the rearchitecture of a monolithic on-premises application into AWS microservices on Docker and Kubernetes, and owned custom haptic VR training systems for enterprise clients end-to-end as primary technical decision-maker.

Japan Exchange Teaching (JET) Program

Tokyo, Japan

ENGLISH LANGUAGE TEACHER

August 2018 - August 2019

- Taught over 1,000 elementary and 500 middle school students, adapting curriculum using self-taught Japanese (JLPT N2).

Certifications & Education

2026 **AWS Certified Machine Learning Engineer - Associate**, AWS Training and Certification

Toronto, Canada

2025 **AWS Certified AI Practitioner**, AWS Training and Certification

Toronto, Canada

2023 **AWS Certified Solutions Architect - Professional**, AWS Training and Certification

Toronto, Canada

2018 **Bachelor of Honors Science - Kinesiology**, McMaster University

Hamilton, Canada